/working with larger datasets

two-approaches

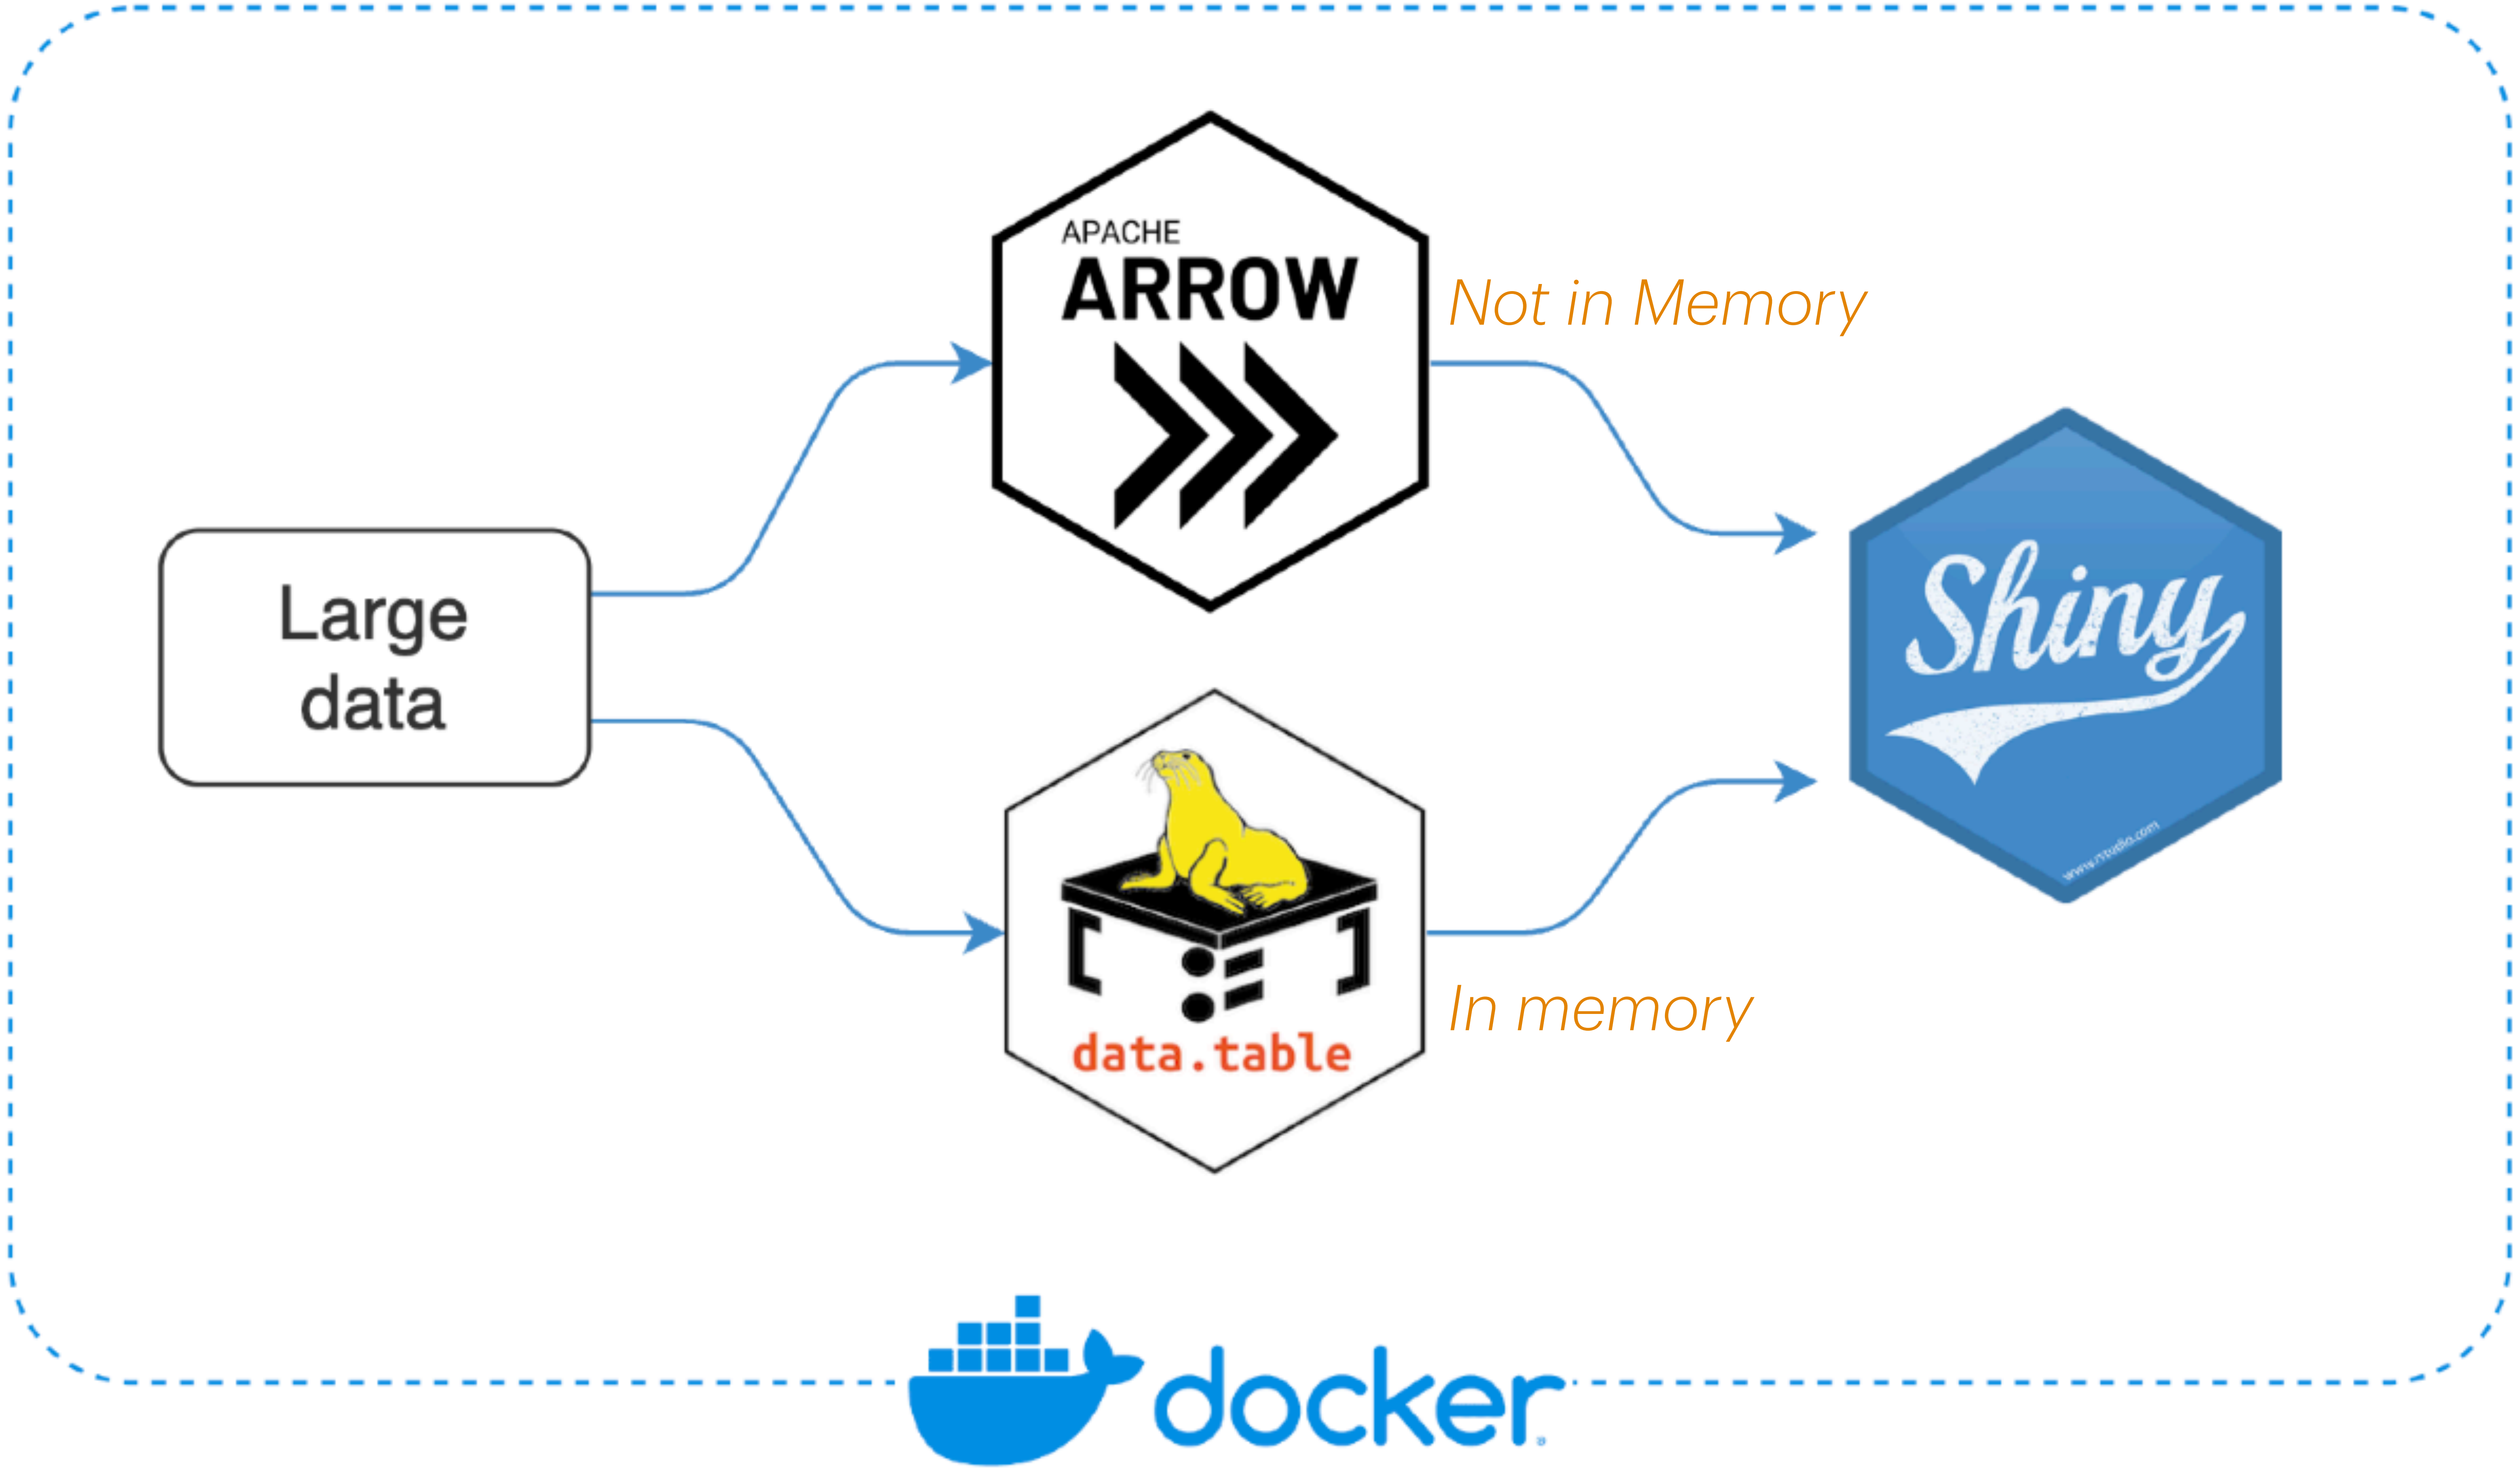

/arrow + parquet

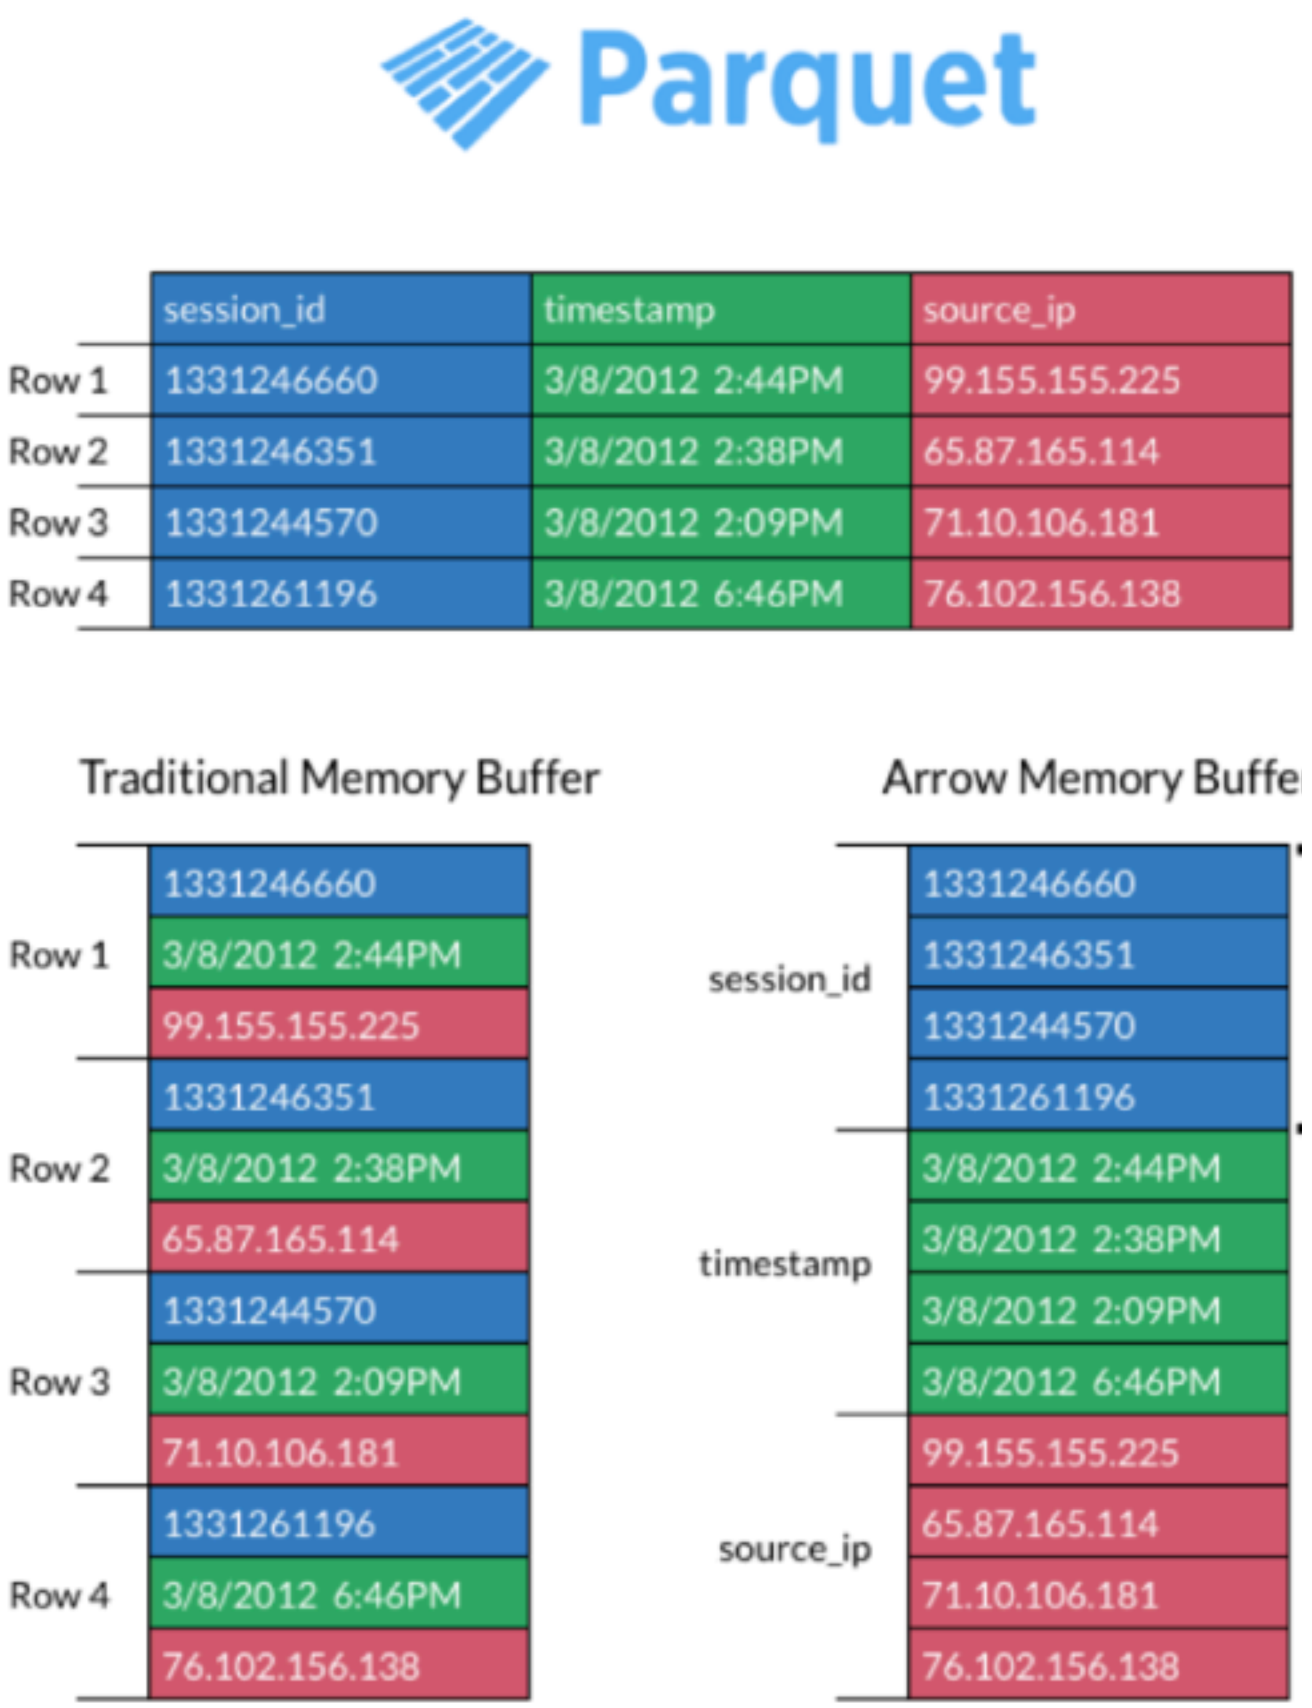

	session_id	timestamp	source_ip
Row 1	1331246660	3/8/2012 2:44PM	99.155.155.225
Row 2	1331246351	3/8/2012 2:38PM	65.87.165.114
Row 3	1331244570	3/8/2012 2:09PM	71.10.106.181
Row 4	1331261196	3/8/2012 6:46PM	76.102.156.138

Traditional Memory Buffer

Row 1	1331246660 3/8/2012 2:44PM 99.155.155.225
Row 2	1331246351 3/8/2012 2:38PM 65.87.165.114
Row 3	1331244570 3/8/2012 2:09PM 71.10.106.181
Row 4	1331261196 3/8/2012 6:46PM 76.102.156.138

Arrow Memory Buffe

session_id	1331246660 1331246351 1331244570 1331261196
timestamp	3/8/2012 2:44PM 3/8/2012 2:38PM 3/8/2012 2:09PM 3/8/2012 6:46PM
source_ip	99.155.155.225 65.87.165.114 71.10.106.181 76.102.156.138

- `{arrow}` evaluates lazily by default
- Verbs: `filter`, `select`, `mutate`, `join`, `distinct`, `group_by` + summarize, and across
- execution only runs on `dplyr::collect()`
- massive performance gains using parquet files and smart partitioning